

## ▼ Linear Models

- In supervised learning, the training data fed to the algorithm includes the desired solutions, called labels.
- In **regression**, the labels are continuous quantities.
- Linear models predict by computing a weighted sum of input features plus a bias term.

```
import numpy as np
%matplotlib inline
import matplotlib
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
np.random.seed(42)
```

5-2

3

## ▼ Simple Linear Regression

Simple linear regression equation:

$$y = ax + b$$

$a$ : slope

$b$ : intercept

Generate linear-looking data with the equation:


$$y = 3X + 4 + noise$$

```
np.random.rand(100, 1)
```

```
array([[0.37454012],
       [0.95071431],
       [0.73199394],
       [0.59865848],
       [0.15601864],
       [0.15599452],
       [0.05808361],
```

```
[0.86617615],  
[0.60111501],  
[0.70807258],  
[0.02058449],  
[0.96990985],  
[0.83244264],  
[0.21233911],  
[0.18182497],  
[0.18340451],  
[0.30424224],  
[0.52475643],  
[0.43194502],  
[0.29122914],  
[0.61185289],  
[0.13949386],  
[0.29214465],  
[0.36636184],  
[0.45606998],  
[0.78517596],  
[0.19967378],  
[0.51423444],  
[0.59241457],  
[0.04645041],  
[0.60754485],  
[0.17052412],  
[0.06505159],  
[0.94888554],  
[0.96563203],  
[0.80839735],  
[0.30461377],  
[0.09767211],  
[0.68423303],  
[0.44015249],  
[0.12203823],  
[0.49517691],  
[0.03438852],  
[0.9093204 ],  
[0.25877998],  
[0.66252228],  
[0.31171108],  
[0.52006802],  
[0.54671028],  
[0.18485446],  
[0.96958463],  
[0.77513282],  
[0.93949894],  
[0.89482735],  
[0.59789998],  
[0.92187424],  
[0.0884925 ],  
[0.10500000]
```

```
X = 2*np.random.rand(100, 1)  
y = 4 + 3 * X + np.random.randn(100, 1)  
plt.scatter(X, y);
```



```
import pandas as pd
pd.DataFrame(y)
```

	0
0	3.508550
1	8.050716
2	6.179208
3	6.337073
4	11.311173
...	...
95	5.441928
96	10.121188
97	9.787643
98	8.061635
99	9.597115

100 rows × 1 columns

```
from sklearn.linear_model import LinearRegression

linear_reg = LinearRegression(fit_intercept=True)
linear_reg.fit(X, y)


LinearRegression()
```

Plot the model's predictions:

```
#X_fit[]

# construct best fit line
X_fit = np.linspace(0, 2, 100)
y_fit = linear_reg.predict(X_fit[:, np.newaxis])

plt.scatter(X, y)
plt.plot(X_fit, y_fit, "r-", linewidth=2, label="Predictions")
plt.xlabel("$X$", fontsize=18)
plt.ylabel("$y$", rotation=0, fontsize=18)
plt.legend(loc="upper left", fontsize=14);
```



Predictions are a good fit.

Generate new data to make predictions with the model:

```
X_new = np.array([[0], [2]])
X_new

array([[0],
       [2]])

X_new.shape

(2, 1)

y_new = linear_reg.predict(X_new)
y_new

array([[ 3.74406122],
       [10.47517611]])
```

```
linear_reg.coef_, linear_reg.intercept_

(array([[3.36555744]]), array([3.74406122]))
```

The model estimates:

$$\hat{y} = 3.36X + 3.74$$

```
#| VENTAS | GANANCIAS |
#COEF*VENTAS+B
#| VENTAS | COMPRAS | GANANCIAS |
#COEF1*X1+COEF2*X2+B=Y
```

## ▼ Polynomial Regression

If data is more complex than a straight line, you can use a linear model to fit non-linear data adding powers of each feature as new features and then train a linear model on the extended set of features.

$$y = a_0 + a_1x_1 + a_2x_2 + a_3x_3 + \dots$$

to

$$y = a_0 + a_1x + a_2x^2 + a_3x^3 + \dots$$


This is still a linear model, the linearity refers to the fact that the coefficients never multiply or divide each other.

To generate polynomial data we use the function:

$$y = 0.50X^2 + X + 2 + noise$$

```
# generate non-linear data e.g. quadratic equation
m = 100
X = 6 * np.random.rand(m, 1) - 3
y = 0.5 * X**2 + X + 2 + np.random.randn(m, 1)

plt.plot(X, y, "b.")
plt.xlabel("$x_1$", fontsize=18)
plt.ylabel("$y$", rotation=0, fontsize=18)
plt.axis([-3, 3, 0, 10]);
```



```
import pandas as pd
pd.DataFrame(y)
```

	0
0	8.529240
1	3.768929
2	3.354423
3	2.747935
4	0.808458
...	...
95	5.346771
96	6.338229
97	3.488785
98	1.372002
99	-0.072150

100 rows × 1 columns

Now we can use `PolynomialFeatures` to transform training data adding the square of each feature as new features.

```
from sklearn.preprocessing import PolynomialFeatures

poly_features = PolynomialFeatures(degree=2, include_bias=False)
X_poly = poly_features.fit_transform(X)
```

`X_poly`

```
array([[ 2.72919168e+00,  7.44848725e+00],
```

```
[ 1.42738150e+00, 2.03741795e+00],  
[ 3.26124315e-01, 1.06357069e-01],  
[ 6.70324477e-01, 4.49334905e-01],  
[-4.82399625e-01, 2.32709399e-01],  
[-1.51361406e+00, 2.29102753e+00],  
[-8.64163928e-01, 7.46779295e-01],  
[ 1.54707666e+00, 2.39344620e+00],  
[-2.91363907e+00, 8.48929262e+00],  
[-2.30356416e+00, 5.30640783e+00],  
[-2.72398415e+00, 7.42008964e+00],  
[-2.75562719e+00, 7.59348119e+00],  
[ 2.13276350e+00, 4.54868016e+00],  
[ 1.22194716e+00, 1.49315485e+00],  
[-1.54957025e-01, 2.40116797e-02],  
[-2.41299504e+00, 5.82254504e+00],  
[-5.03047493e-02, 2.53056780e-03],  
[-1.59169375e-01, 2.53348900e-02],  
[-1.96078878e+00, 3.84469264e+00],  
[-3.96890105e-01, 1.57521755e-01],  
[-6.08971594e-01, 3.70846402e-01],  
[ 6.95100588e-01, 4.83164828e-01],  
[ 8.10561905e-01, 6.57010602e-01],  
[-2.72817594e+00, 7.44294397e+00],  
[-7.52324312e-01, 5.65991871e-01],  
[ 7.55159494e-01, 5.70265862e-01],  
[ 1.88175515e-02, 3.54100244e-04],  
[ 2.13893905e+00, 4.57506025e+00],  
[ 9.52161790e-01, 9.06612074e-01],  
[-2.02239344e+00, 4.09007522e+00],  
[-2.57658752e+00, 6.63880323e+00],  
[ 8.54515669e-01, 7.30197029e-01],  
[-2.84093214e+00, 8.07089541e+00],  
[ 5.14653488e-01, 2.64868212e-01],  
[ 2.64138145e+00, 6.97689596e+00],  
[ 4.52845067e-01, 2.05068655e-01],  
[-6.70980443e-01, 4.50214755e-01],  
[ 8.59729311e-01, 7.39134488e-01],  
[-2.50482657e-01, 6.27415615e-02],  
[ 2.73700736e-01, 7.49120928e-02],  
[ 2.64878885e+00, 7.01608239e+00],  
[-6.83384173e-01, 4.67013928e-01],  
[ 2.76714338e+00, 7.65708250e+00],  
[ 2.43210385e+00, 5.91512915e+00],  
[-1.82525319e+00, 3.33154921e+00],  
[-2.58383219e+00, 6.67618881e+00],  
[-2.39533199e+00, 5.73761535e+00],  
[-2.89066905e+00, 8.35596753e+00],  
[-2.43334224e+00, 5.92115443e+00],  
[ 1.09804064e+00, 1.20569325e+00],  
[-2.57286811e+00, 6.61965031e+00],  
[-1.08614622e+00, 1.17971361e+00],  
[ 2.06925187e+00, 4.28180328e+00],  
[-2.86036839e+00, 8.18170730e+00],  
[ 1.88681090e+00, 3.56005536e+00],  
[-1.30887135e+00, 1.71314421e+00],  
[-2.29101103e+00, 5.24873156e+00],
```

```

[ 1.18042200e+00  1.20220844e+00]

```

`X_poly` now contains the original feature of `X` plus the square of the feature:

```

print(X[0])
print(X[0]*X[0])

```

```

[2.72919168]
[7.44848725]

```

```

X_poly[0]

array([2.72919168, 7.44848725])

```

Fit the model to this extended training data:

```

lin_reg = LinearRegression(fit_intercept=True)
lin_reg.fit(X_poly, y)
lin_reg.coef_, lin_reg.intercept_

(array([[1.04271531, 0.50866711]]), array([2.01873554]))

```

The model estimates:

$$\hat{y} = 0.89X + 0.48X^2 + 2.09$$


Plot the data and the predictions:

```

X_new=np.linspace(-3, 3, 100).reshape(100, 1)
X_new_poly = poly_features.transform(X_new)
y_new = lin_reg.predict(X_new_poly)
plt.plot(X, y, "b.")
plt.plot(X_new, y_new, "r-", linewidth=2, label="Predictions")
plt.xlabel("$x_1$", fontsize=18)
plt.ylabel("$y$", rotation=0, fontsize=18)
plt.legend(loc="upper left", fontsize=14)
plt.axis([-3, 3, 0, 10]);

```





## R square

$R^2$  es una medida estadística de qué tan cerca están los datos de la línea de regresión ajustada. También se conoce como el coeficiente de determinación o el coeficiente de determinación múltiple para la regresión múltiple. Para decirlo en un lenguaje más simple,  $R^2$  es una medida de ajuste para los modelos de regresión lineal.

$R^2$  no indica si un modelo de regresión se ajusta adecuadamente a sus datos. Un buen modelo puede tener un valor  $R^2$  bajo. Por otro lado, un modelo sesgado puede tener un valor alto de  $R^2$ .

$SS_{res} + SS_{reg} = SS_{tot}$ ,  $R^2 = \text{Explained variation} / \text{Total Variation}$

$$R^2 = 1 - \frac{\text{Sum Squared Regression Error } SS_{Regression}}{\text{Sum Squared Total Error } SS_{Total}}$$

$$R^2 \equiv 1 - \frac{SS_{res}}{SS_{tot}} \equiv 1 - \frac{\sum (y_i - \hat{y}_i)^2}{\sum (y_i - \bar{y})^2}$$

$$\downarrow$$

$$R^2 = \frac{SS_{reg}}{SS_{tot}}$$

## ▼ Ejercicio 1

Utiliza la base de datos de <https://www.kaggle.com/vinicius150987/manufacturing-cost>

Suponga que trabaja como consultor de una empresa de nueva creación que busca desarrollar un modelo para estimar el costo de los bienes vendidos a medida que varían el volumen de producción (número de unidades producidas). La startup recopiló datos y le pidió que desarrollara un modelo para predecir su costo frente a la cantidad de unidades vendidas.

```
import pandas as pd
df = pd.read_csv('https://raw.githubusercontent.com/marypazrf/bdd/main/EconomiesOfSca
df.sample(10)
```

	Number of Units	Manufacturing Cost
968	7.065653	27.804027
212	3.372115	41.127212
416	4.194513	43.832711
677	5.068888	41.225741
550	4.604122	37.569764
764	5.389522	31.191501
386	4.104190	42.988730
339	3.942214	46.291435
82	2.665856	48.578425
487	4.399514	37.567914

```
X = df[['Number of Units']]
y = df['Manufacturing Cost']
```

```
len(X)
```

```
1000
```

```
y.describe
```

```
<bound method NDFrame.describe of 0      95.066056
1      96.531750
2      73.661311
3      95.566843
4      98.777013
...
995    23.855067
996    27.536542
```

```
997    25.973787
```


```
998    25.138311
```

```
999    21.547777
```

```
Name: Manufacturing Cost, Length: 1000, dtype: float64>
```

```
plt.plot(X,y,'b.')
```

```
[<matplotlib.lines.Line2D at 0x7f72cce23110>]
```



```
#lineal
```

```
linear_reg = LinearRegression(fit_intercept=True)
```

```
linear_reg.fit(X,y)
```

```
LinearRegression()
```

```
linear_reg.coef_
```

```
array([-6.0333683])
```

```
linear_reg.intercept_
```

```
67.03904193389238
```

```
y' = -6.03X + 67.03
```

```
y_pred = linear_reg.predict(X)
```

```
from sklearn import metrics
```

```
from sklearn.metrics import r2_score
```

```
print('Error medio Absoluto (MAE):', metrics.mean_absolute_error(y, y_pred))
```

```
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y, y_pred)))
```


```
print('r2_score',r2_score(y, y_pred))
```

Error medio Absoluto (MAE): 4.921245677483368  
 Root Mean Squared Error: 6.871474273023592  
 r2\_score 0.5789767558506287

```
X_fit = np.linspace(0, 10, 100)
y_fit = linear_reg.predict(X_fit[:, np.newaxis])
```

```
plt.scatter(X, y)
plt.plot(X_fit, y_fit, "r-", linewidth=2, label="Predictions")
plt.xlabel("$X$", fontsize=18)
plt.ylabel("$y$", rotation=0, fontsize=18)
plt.legend(loc="upper left", fontsize=14);
```

/usr/local/lib/python3.7/dist-packages/sklearn/base.py:451: UserWarning: X does  
 "X does not have valid feature names, but"



```
#polinomial
poly_transform = PolynomialFeatures(degree=2, include_bias = False)
polyX = poly_transform.fit_transform(X)
```

```
polyX
```

```
array([[ 1.          ,  1.          ],
       [ 1.18599365,  1.40658094],
       [ 1.19149864,  1.41966901],
       ...,
       [ 8.78088812,  77.10399623],
       [ 8.89769971,  79.16906006],
       [10.          , 100.          ]])
```

```
poly_regression = LinearRegression(fit_intercept=True)
poly_regression.fit(polyX, y)
```

```
LinearRegression()
```

```
poly_regression.coef_
```

```
array([-16.82425278,  1.17336718])
```

```
poly_regression.intercept_
```


```
89.73719677939121
```

```
y' = -16.82X + 1.17x^2 + 89.73
```

```
y_pred = poly_regression.predict(polyX)
```

```
X_fit = np.linspace(0, 10, 100).reshape(100,1)
polyX_fit = poly_transform.fit_transform(X_fit)
y_fit = poly_regression.predict(polyX_fit)
```

```
plt.scatter(X, y)
plt.plot(X_fit, y_fit, "r-", linewidth=2, label="Predictions")
plt.xlabel("$X$", fontsize=18)
plt.ylabel("$y$", rotation=0, fontsize=18)
plt.legend(loc="upper left", fontsize=14);
```



```
from sklearn import metrics
from sklearn.metrics import r2_score
```

```
print('Error medio Absoluto (MAE):', metrics.mean_absolute_error(y, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y, y_pred)))
print('r2_score', r2_score(y, y_pred))
```

Error medio Absoluto (MAE): 4.538765565228896  
Root Mean Squared Error: 6.129173450227568  
r2\_score 0.6650268116056028

```
df = pd.read_csv('https://raw.githubusercontent.com/marypazrf/bdd/main/EconomiesOfSca  
df.head()
```

	Number of Units	Manufacturing Cost
0	1.000000	95.066056
1	1.185994	96.531750
2	1.191499	73.661311
3	1.204771	95.566843
4	1.298773	98.777013

```
X = df[['Number of Units']]  
y = df['Manufacturing Cost']
```

```
from sklearn.model_selection import train_test_split  
from sklearn.preprocessing import StandardScaler  
from sklearn.pipeline import Pipeline  
np.random.seed(42)  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_sta
```

Lineal Regression

```
X_train.shape
```

(800, 1)

```
y_train.shape
```


(800,)

```
y_test.shape
```

(200,)

```
plt.scatter(X_train, y_train)
```

<matplotlib.collections.PathCollection at 0x7f72ccd28a90>



```
from sklearn.linear_model import LinearRegression
```


```
linear_reg = LinearRegression(fit_intercept=True)
linear_reg.fit(X_train, y_train)
```

```
LinearRegression()
```

```
X_fit = np.linspace(0, 10, 100)
y_fit = linear_reg.predict(X_fit[:, np.newaxis])
```

```
plt.scatter(X_test, y_test)
plt.plot(X_fit, y_fit, "r-", linewidth=2, label="Predictions")
plt.xlabel("$X$", fontsize=18)
plt.ylabel("$y$", rotation=0, fontsize=18)
plt.legend(loc="upper left", fontsize=14);
```

/usr/local/lib/python3.7/dist-packages/sklearn/base.py:451: UserWarning: X does not have valid feature names, but



```
linear_reg.coef_
```

```
array([-6.22263745])
```

```
linear_reg.intercept_
```

```
68.08187672497846
```

```
y_pred = linear_reg.predict(X_train)
from sklearn import metrics
from sklearn.metrics import r2_score
linear_Mae_train=metrics.mean_absolute_error(y_train,y_pred)
linear_r2_train=r2_score(y_train,y_pred)
print('Error medio Absoluto (MAE):', metrics.mean_absolute_error(y_train, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_train, y_pred))
print('r2_score', r2_score(y_train, y_pred))
```

```
Error medio Absoluto (MAE): 4.983664639039784
Root Mean Squared Error: 6.930601695798021
r2_score 0.594243664202029
```

```
y_pred = linear_reg.predict(X_test)
from sklearn import metrics
from sklearn.metrics import r2_score
linear_Mae=metrics.mean_absolute_error(y_test,y_pred)
linear_r2=r2_score(y_test,y_pred)
print('Error medio Absoluto (MAE):', metrics.mean_absolute_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred))
print('r2_score', r2_score(y_test, y_pred))
```

```
Error medio Absoluto (MAE): 4.83324114398872
Root Mean Squared Error: 6.6682038827508405
r2_score 0.48439969577999376
```

## Polynomial Regression

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_sta
```

```
y_train.shape
```


```
(800,)
```

```
y_test.shape
```



(200,)

```
plt.plot(X_train, y_train, "b.")
plt.xlabel("$x_1$", fontsize=18)
plt.ylabel("$y$", rotation=0, fontsize=18)
plt.show()
```



```
from sklearn.preprocessing import PolynomialFeatures
```

```
poly_transform = PolynomialFeatures(degree=3, include_bias = False)
polyX = poly_transform.fit_transform(X_train)
polyX
```

```
array([[ 5.65669179,  31.99816198, 181.00374013],
       [ 5.53635996,  30.65128163, 169.69652839],
       [ 2.25713058,   5.09463847,  11.49926431],
       ...,
       [ 2.64655085,   7.00423138,  18.53705446],
       [ 5.18000789,  26.83248172, 138.99246696],
       [ 6.37652342,  40.66005096, 259.26976735]])
```

```
poly_regression = LinearRegression(fit_intercept=True)
poly_regression.fit(polyX, y_train)
```

```
LinearRegression()
```

```
poly_regression.coef_
```

```
array([-35.3746951 ,   5.28612952,  -0.28112892])
```

```
poly_regression.intercept_
```

115.22180379591046

```
poly_transform = PolynomialFeatures(degree=3, include_bias = False)
polyX = poly_transform.fit_transform(X_train)
```

```
y_pred = poly_regression.predict(polyX)
```


```
from sklearn import metrics
from sklearn.metrics import r2_score
Poly_Mae_train=metrics.mean_absolute_error(y_train,y_pred)
Poly_r2_train=r2_score(y_train,y_pred)
print('Error medio Absoluto (MAE):', metrics.mean_absolute_error(y_train, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_train, y_pred))
print('r2_score',r2_score(y_train, y_pred ))
```

```
Error medio Absoluto (MAE): 4.462630935588183
Root Mean Squared Error: 5.858467673318268
r2_score 0.7100711021026239
```

```
X_fit = np.linspace(0, 10, 100).reshape(100,1)
polyX_fit = poly_transform.fit_transform(X_fit)
```

```
y_fit = poly_regression.predict(polyX_fit)
```

```
plt.scatter(X_test, y_test)
plt.plot(X_fit, y_fit, "r-", linewidth=2, label="Predictions")
plt.xlabel("$X$", fontsize=18)
plt.ylabel("$y$", rotation=0, fontsize=18)
plt.legend(loc="upper left", fontsize=14);
```



```
poly_transform = PolynomialFeatures(degree=3, include_bias = False)
polyX = poly_transform.fit_transform(X_test)
```

```
y_pred = poly_regression.predict(polyX)
```

```
from sklearn import metrics
from sklearn.metrics import r2_score
Poly_Mae=metrics.mean_absolute_error(y_test,y_pred)
Poly_r2=r2_score(y_test,y_pred)
print('Error medio Absoluto (MAE):', metrics.mean_absolute_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
print('r2_score',r2_score(y_test, y_pred ))
```

```
Error medio Absoluto (MAE): 4.6058129615113925
```


```
Root Mean Squared Error: 6.209999522706177
```

```
r2_score 0.5528239263527293
```


## Ridge

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_sta
```


```
plt.plot(X, y,"b.")
plt.xlabel("$X$", fontsize=15)
plt.ylabel("$y$", rotation=0, fontsize=15);
```



```
plt.plot(X_train, y_train,"b.")
plt.xlabel("$X$", fontsize=15)
plt.ylabel("$y$", rotation=0, fontsize=15);
```



```
plt.plot(X_test, y_test,"b.")
plt.xlabel("$X$", fontsize=18)
plt.ylabel("$y$", rotation=0, fontsize=18);
```



```
X_train2=X_train.to_numpy().reshape(-1,1)
```


```
model = Pipeline([("poly_features", PolynomialFeatures(degree=2, include_bias=False))
                  ("scaler", StandardScaler()),
                  ("linear_reg", LinearRegression())])
model.fit(X_train, y_train)
```

```
Pipeline(steps=[('poly_features', PolynomialFeatures(include_bias=False)),
                 ('scaler', StandardScaler()),
                 ('linear_reg', LinearRegression())])
```

```
y_pred = model.predict(X_test)
```

```
order = np.argsort(X_test.values.ravel())
X_test2 = X_test.values.ravel()[order]
y_test2 = y_test.values[order]
y_pred2 = y_pred[order]
```

```
plt.plot(X, y, "b.")
plt.plot(X_test2, y_pred2, "r-", linewidth=1, label="No Regularization")
plt.xlabel("$X$", fontsize=18)
plt.ylabel("$y$", rotation=0, fontsize=18)
plt.legend(loc="upper right", fontsize=10);
```



```
from sklearn.linear_model import Ridge
```

```
model_ridge = Pipeline([("poly_features", PolynomialFeatures(degree=4, include_bias=False),
    ("scaler", StandardScaler()),
    ("ridge", Ridge(alpha=1, solver = 'cholesky', random_state = 42,max
model_ridge.fit(X_train, y_train)
```

```
Pipeline(steps=[('poly_features',
    PolynomialFeatures(degree=4, include_bias=False)),
    ('scaler', StandardScaler()),
    ('ridge',
    Ridge(alpha=1, max_iter=1000, random_state=42,
    solver='cholesky'))])
```

```
y_pred = model_ridge.predict(X_train)
Ridge_Mae_train=metrics.mean_absolute_error(y_train,y_pred)
Ridge_r2_train=r2_score(y_train,y_pred)
print('Error medio Absoluto (MAE):', metrics.mean_absolute_error(y_train, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_train, y_pred))
print('r2_score', r2_score(y_train, y_pred))
```


```
Error medio Absoluto (MAE): 4.453130806314608
Root Mean Squared Error: 5.948178018959322
r2_score 0.7011237920037103
```

```
y_pred = model_ridge.predict(X_test)
```

```
order = np.argsort(X_test.values.ravel())
```

```
X_test2 = X_test.values.ravel()[order]
y_test2 = y_test.values[order]
y_pred2 = y_pred[order]
```

```
plt.plot(X_test, y_test, "b.")
plt.plot(X_test2, y_pred2, "g-", label="with regularization")
plt.xlabel("$X$", fontsize=12)
plt.ylabel("$y$", rotation=0, fontsize=12)
plt.legend(loc="upper right", fontsize=10);
```



```
X_test=X_test.to_numpy().reshape(-1,1)
y_pred = model.predict(X_test)
from sklearn import metrics
from sklearn.metrics import r2_score
Ridge_Mae=metrics.mean_absolute_error(y_test,y_pred)
Ridge_r2=r2_score(y_test,y_pred)
print('Error medio Absoluto (MAE):', metrics.mean_absolute_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
print('r2_score',r2_score(y_test, y_pred ))
```

Error medio Absoluto (MAE): 4.762293483795133

Root Mean Squared Error: 6.513768934466413

r2\_score 0.508005642552175

/usr/local/lib/python3.7/dist-packages/sklearn/base.py:451: UserWarning: X does  
"X does not have valid feature names, but"

lasso

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_sta
```

```
from sklearn.linear_model import Lasso
```

```
model_lasso = Pipeline([("poly_features", PolynomialFeatures(degree=15, include_bias=
    ("scaler", StandardScaler()),
    ("lasso", Lasso(alpha = .9, random_state = 10, tol=0.1, max_iter
model_lasso.fit(X_train, y_train)
```

```
    Pipeline(steps=[('poly_features',
        PolynomialFeatures(degree=15, include_bias=False)),
        ('scaler', StandardScaler()),
        ('lasso', Lasso(alpha=0.9, random_state=10, tol=0.1))])
```

```
y_pred = model_lasso.predict(X_train)
```

```
Lasso_MAE_train=metrics.mean_absolute_error(y_train,y_pred)
Lasso_r2_train=r2_score(y_train,y_pred)
```

```
print('Error medio Absoluto (MAE):', metrics.mean_absolute_error(y_train, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_train, y_pred))
print('r2_score',r2_score(y_train, y_pred ))
```

```
Error medio Absoluto (MAE): 4.8969775859486955
Root Mean Squared Error: 6.935898287947112
r2_score 0.5936232427341488
```

```
y_pred = model_lasso.predict(X_test)
```

```
order = np.argsort(X_test.values.ravel())
X_test2 = X_test.values.ravel()[order]
y_test2 = y_test.values[order]
y_pred2 = y_pred[order]
```

```
plt.plot(X_test, y_test, "b.")
plt.plot(X_test2, y_pred2, "g-",label="Predictions Lasso Regularization")
plt.xlabel("$X$", fontsize=18)
plt.ylabel("$y$", rotation=0, fontsize=18)
plt.legend(loc="upper right", fontsize=10);
```



```
y_pred = model_lasso.predict(X_test)
from sklearn import metrics
from sklearn.metrics import r2_score
Lasso_MAE=metrics.mean_absolute_error(y_test,y_pred)
Lasso_r2=r2_score(y_test,y_pred)
print('Error medio Absoluto (MAE):', metrics.mean_absolute_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
print('r2_score',r2_score(y_test, y_pred ))
```

```
Error medio Absoluto (MAE): 4.646351122634725
Root Mean Squared Error: 6.486538755327113
r2_score 0.51211051393966
```

## MAE

```
def diagramas(metrics):
    sumt=list()
    diagrama=metrics
    sumt.append(diagrama)
    return sumt
```

```
Metricas = list()
Metricas.append('Lin train')
Metricas.append('Lin')
Metricas.append('Poly train')
Metricas.append('Poly')
Metricas.append('Ridge train')
Metricas.append('Ridge')
Metricas.append('Lasso train')
Metricas.append('Lasso')
```

```
MAE_list = list()
MAE_list.append(linear_Mae_train)
MAE_list.append(linear_Mae)
MAE_list.append(Poly_Mae_train)
MAE_list.append(Poly_Mae)
MAE_list.append(Ridge_Mae_train)
MAE_list.append(Ridge_Mae)
MAE_list.append(Lasso_MAE_train)
MAE_list.append(Lasso_MAE)
```

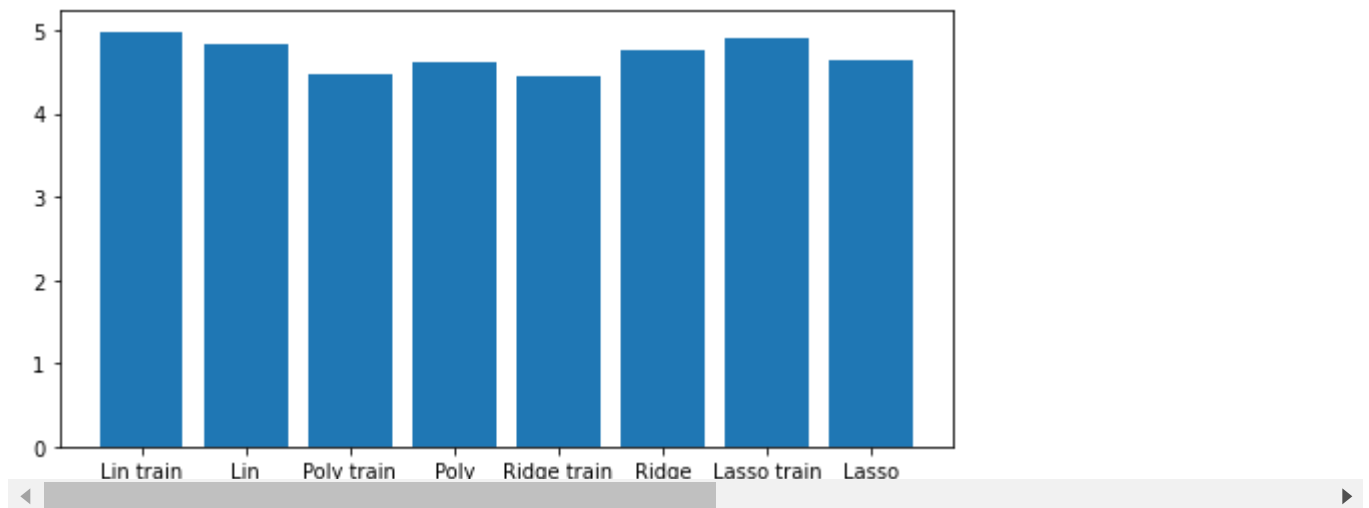


```
fig, ax = plt.subplots(figsize =(8, 4))
print("MAE")
print(MAE_list)
Metricas=list(Metricas)
plt.bar(Metricas,MAE_list)
```

```
plt.show()
```

MAE


[4.983664639039784, 4.83324114398872, 4.462630935588183, 4.6058129615113925, 4.4



```
R2_list = list()
R2_list.append(linear_r2_train)
R2_list.append(linear_r2)
R2_list.append(Poly_r2_train)
R2_list.append(Poly_r2)
R2_list.append(Ridge_r2_train)
R2_list.append(Ridge_r2)
R2_list.append(Lasso_r2_train)
R2_list.append(Lasso_r2)
```

```
fig, ax = plt.subplots(figsize =(8, 4))
print("r2")
print(R2_list)
R2=list(R2_list)
plt.bar(Metricas,R2_list)
```

```
plt.show()
```



## Conclusiones del ejercicio 1

La regression lineal no fue buena para los datos que se tienen. Mostro mejor comportamiento el modelo polinomial.

Para la empresa convendria usar el modelo MAE pero se requiere hacer mas pruebas.

## ▼ Ejercicio 2

Realiza la regresión polinomial de los siguientes datos:

```
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import Lasso
from sklearn.linear_model import Ridge
from sklearn import metrics
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.metrics import r2_score
import numpy as np
%matplotlib inline
import matplotlib
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
np.random.seed(42)
df = pd.read_csv('https://raw.githubusercontent.com/marypazrf/bdd/main/kc_house_data.
df.sample(10)
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqf
<b>735</b>	2591820310	20141006T000000	365000.0	4	2.25	2070	
<b>2830</b>	7974200820	20140821T000000	865000.0	5	3.00	2900	
<b>4106</b>	7701450110	20140815T000000	1038000.0	4	2.50	3770	
<b>16218</b>	9522300010	20150331T000000	1490000.0	3	3.50	4560	
<b>19964</b>	9510861140	20140714T000000	711000.0	3	2.50	2550	
<b>1227</b>	1761300310	20140827T000000	211000.0	4	2.00	1710	
<b>18849</b>	7732410120	20140819T000000	790000.0	4	2.50	2690	
<b>19369</b>	7010701383	20141017T000000	680000.0	3	2.50	1800	
<b>20164</b>	291310170	20140804T000000	384500.0	3	2.50	1600	
<b>7139</b>	4232901990	20140516T000000	605000.0	2	1.00	910	

10 rows × 21 columns

df.info()


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 21 columns):
#   Column              Non-Null Count  Dtype
---  -
0   id                   21613 non-null  int64
1   date                 21613 non-null  object
2   price                21613 non-null  float64
3   bedrooms             21613 non-null  int64
4   bathrooms            21613 non-null  float64
5   sqft_living          21613 non-null  int64
6   sqft_lot             21613 non-null  int64
7   floors               21613 non-null  float64
8   waterfront           21613 non-null  int64
9   view                 21613 non-null  int64
10  condition            21613 non-null  int64
11  grade                21613 non-null  int64
12  sqft_above           21613 non-null  int64
13  sqft_basement        21613 non-null  int64
14  yr_built              21613 non-null  int64
15  yr_renovated         21613 non-null  int64
16  zipcode              21613 non-null  int64
17  lat                  21613 non-null  float64
18  long                 21613 non-null  float64
19  sqft_living15        21613 non-null  int64
20  sqft_lot15           21613 non-null  int64
dtypes: float64(5), int64(15), object(1)
memory usage: 3.5+ MB
```

df.describe()

	<b>id</b>	<b>price</b>	<b>bedrooms</b>	<b>bathrooms</b>	<b>sqft_living</b>	<b>sqft_lot</b>
<b>count</b>	2.161300e+04	2.161300e+04	21613.000000	21613.000000	21613.000000	2.161300e+04
<b>mean</b>	4.580302e+09	5.400881e+05	3.370842	2.114757	2079.899736	1.510697e+04
<b>std</b>	2.876566e+09	3.671272e+05	0.930062	0.770163	918.440897	4.142051e+04
<b>min</b>	1.000102e+06	7.500000e+04	0.000000	0.000000	290.000000	5.200000e+02
<b>25%</b>	2.123049e+09	3.219500e+05	3.000000	1.750000	1427.000000	5.040000e+03
<b>50%</b>	3.904930e+09	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03
<b>75%</b>	7.308900e+09	6.450000e+05	4.000000	2.500000	2550.000000	1.068800e+04
<b>max</b>	9.900000e+09	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06

```
df.drop('id', axis = 1, inplace = True)
df.drop('date', axis = 1, inplace = True)
df.drop('zipcode', axis = 1, inplace = True)
df.drop('lat', axis = 1, inplace = True)
df.drop('long', axis = 1, inplace = True)
```

```
plt.figure(figsize=(12,8))
sns.heatmap(df.corr(), annot=True, cmap='Dark2_r', linewidths = 2)
plt.show()
```



```
columns = df.columns.drop('price')
```

```
features = columns
```

```
label = ['price']
```

```
X = df[features]
```

```
y = df[label]
```

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.1, random_state=42)
```

```
print(f'Numero total de registros en la bdd: {len(X)}')
```

```
print("*****" * 10)
```

```
print(f'Numero total de registros en el training set: {len(X_train)}')
```

```
print(f'Tamaño de X_train: {X_train.shape}')
```

```
print("*****" * 10)
```

```
print(f'Numero total de registros en el test dataset: {len(X_test)}')
```

```
print(f'Tamaño del X_test: {X_test.shape}')
```

```
Numero total de registros en la bdd: 21613
```

```
*****
```

```
Numero total de registros en el training set: 19451
```

```
Tamaño de X_train: (19451, 15)
```

```
*****
```

```
Mumero total de registros en el test dataset: 2162
```

```
Tamaño del X_test: (2162, 15)
```

```
#tu codigo aquí
```

```
plt.plot(X_train,y_train,'b.')
```

```
plt.show()
```



```
linear_reg = LinearRegression(fit_intercept=True)
linear_reg.fit(X_train, y_train)
```

```
LinearRegression()
```

```
. | 2.83 |
```

```
linear_reg.coef_,
```

```
(array([[ -3.82008048e+04,  4.14661380e+04,  1.07992584e+02,
          1.71356997e-02,  3.16916913e+04,  5.52691023e+05,
          4.12493228e+04,  2.12221443e+04,  1.19493216e+05,
          4.77750271e+01,  6.02175565e+01, -3.55090216e+03,
          1.32602215e+01,  2.90059284e+01, -5.48132603e-01]]),)
```

```
linear_reg.intercept_
```

```
array([6151359.2627409])
```

```
y_pred = linear_reg.predict(X_train)
from sklearn import metrics
from sklearn.metrics import r2_score
linear_Mae_train=metrics.mean_absolute_error(y_train,y_pred)
linear_r2_train=r2_score(y_train,y_pred)
print('Error medio Absoluto (MAE):', metrics.mean_absolute_error(y_train, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_train, y_pred))
print('r2_score', r2_score(y_train, y_pred))
```

```
Error medio Absoluto (MAE): 139269.32939115583
Root Mean Squared Error: 214234.8822754647
r2_score 0.6529196653133833
```

```
y_pred = linear_reg.predict(X_test)
from sklearn import metrics
from sklearn.metrics import r2_score
linear_Mae=metrics.mean_absolute_error(y_test,y_pred)
linear_r2=r2_score(y_test,y_pred)
print('Error medio Absoluto (MAE):', metrics.mean_absolute_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred))
print('r2_score', r2_score(y_test, y_pred))
```

```
Error medio Absoluto (MAE): 137480.1388273178
Root Mean Squared Error: 232133.3676240749
r2_score 0.6579723205007814
```

```
from sklearn.preprocessing import PolynomialFeatures
```

```
poly_transform = PolynomialFeatures(degree=3, include_bias = False)
polyX = poly_transform.fit_transform(X_train)
polyX
```

```
array([[5.00000000e+00, 2.75000000e+00, 3.75000000e+03, ...,
        7.06230720e+10, 1.64221278e+11, 3.81867106e+11],
       [4.00000000e+00, 4.50000000e+00, 5.25000000e+03, ...,
        3.02200000e+11, 5.84478976e+12, 1.13042910e+14],
       [3.00000000e+00, 2.50000000e+00, 2.88000000e+03, ...,
        9.21194624e+10, 3.27210820e+11, 1.16226168e+12],
       ...,
       [3.00000000e+00, 2.25000000e+00, 1.78000000e+03, ...,
        1.73857625e+10, 8.77923438e+10, 4.43322266e+11],
       [2.00000000e+00, 1.00000000e+00, 1.15000000e+03, ...,
        9.24639408e+09, 2.73030912e+10, 8.06215680e+10],
       [3.00000000e+00, 1.00000000e+00, 1.45000000e+03, ...,
        8.37158400e+09, 6.23039040e+10, 4.63684824e+11]])
```

```
poly_regression = LinearRegression(fit_intercept=True)
poly_regression.fit(polyX, y_train)
poly_regression.intercept_
```

```
array([19759145.24258333])
```

```
poly_regression.coef_
```

```
array([[ 3.17395792e+02, -2.70130239e+02,  1.30610945e+01,
        -1.40757708e+01,  2.62496809e+00,  3.24633595e+00,
        -8.08938960e-01,  3.60826046e-01, -6.17577143e-01,
         1.54124566e-01,  9.07922470e-01, -2.42302232e-01,
         1.57944629e-01, -4.12303248e-01, -1.89494993e+01,
         2.57497810e-01,  3.31180922e-01,  2.56156276e+01,
        -9.51793255e-02, -6.36531339e-01,  6.24165339e-01,
        -6.14143248e-02,  9.57006257e-01,  1.21507890e+00,
         1.85183285e+01,  7.24785462e+00,  3.94289500e+02,
         4.24069045e+00,  5.07845227e+00, -1.89095034e+01,
        -4.03258004e-01, -2.08156624e+02, -6.15751437e+01,
        -4.82898236e-01, -6.40532068e-02, -8.73845428e-02,
        -1.73829295e-01, -5.97615346e-01, -1.69018199e+02,
        -3.91260224e+01, -1.76779529e+02, -4.47565136e+00,
        -1.70541213e+02,  1.09037853e+02, -3.48181375e-01,
         5.54961266e-03, -3.10542727e+02, -2.04577118e+01,
        -5.10909707e+01,  1.44081886e+02,  1.11581166e+02,
         2.16670060e-01, -5.64909219e-01, -1.27410167e-01,
        -1.12568260e-01,  1.58018890e+00, -3.54328039e-02,
         4.38742488e-05,  6.16058584e+01, -2.69898947e+01,
         1.91781121e+01,  1.20426825e+01,  1.31650265e+01,
         9.34375304e-03, -3.73525411e-03, -1.04398912e-01,
         1.38884709e-02,  3.01381870e-02,  6.78520767e-05,
        -5.34102837e-01, -5.78470986e-02, -1.13822284e-01,
```

```
-4.03570695e-01, -1.05966620e+00, -2.67524020e+02,
-4.30187377e+01, -2.93222557e+02, -3.35793147e+00,
-2.91958127e+02, -9.76415380e+01, -1.57122761e-02,
-1.68119414e-01, -1.13491517e-01, -1.64619588e-01,
-1.74203932e+01, -3.03731175e+00, -1.52993218e+01,
2.89521925e-01, -1.93075606e+01, -1.08573905e+02,
-1.08424368e-01, -7.86651310e-02, -1.31888702e-01,
-4.03527763e+01, -1.07382336e+01, -4.51049163e+01,
-3.19195216e+00, -5.69571664e+01, -2.42072671e+01,
5.39400863e-01, 6.90456940e-01, 1.18106862e+02,
2.59750399e+01, 1.42884610e+02, -1.84020119e+00,
1.37140757e+02, 4.79912276e+00, 7.61197307e-01,
8.82755148e+01, 2.33056421e+01, 9.42517326e+01,
7.84046146e-01, 9.95172836e+01, -1.51103170e+01,
-2.02758252e+00, 2.24427261e+00, 9.40121740e-01,
-1.85322322e+00, 1.45017708e+00, -1.97077053e-03,
-2.80914700e+00, -1.06751495e+00, 1.74066651e+00,
1.30002356e-01, -3.34363391e-02, -1.65748055e+01,
4.33882740e+01, 1.57874743e+00, 2.34932788e-02,
-4.25386312e+01, -4.37936627e-01, 8.20919268e-02,
-1.48254749e+00, 6.97841641e-03, 3.89241075e-04,
1.29106646e+02, -7.16509286e-01, 4.15395793e+00,
2.29637081e-01, 4.75102521e-01, -7.01819595e-01,
-1.06920772e+00, 7.56507260e+00, 2.69567648e+00,
-4.22023433e+00, 8.37419219e+00, -6.85156850e+00,
3.79756814e+00, 3.61539631e-01, 1.60344901e-02,
-1.34799509e+00, 2.13243059e+00, -6.39846626e-01,
-1.28598317e+00, -1.03087982e-01, 1.41973419e-01,
4.08102711e-01, -1.10957527e-01, -8.95127159e+00,
1.10837021e+01, -2.37618609e+01, -1.72495794e+01,
3.64892153e+01, 5.39841200e-01, -5.29031580e-03,
5.17471543e-06, -1.82434474e+01, 2.48592345e+01,
9.70652876e+00, -2.58285326e+00, -1.18894778e+00,
1.25486887e-02, -1.78456991e-02, 1.38371287e-02,
```

```
polyX = poly_transform.fit_transform(X_train)
```

```
y_pred = poly_regression.predict(polyX)
```

```
from sklearn import metrics
```

```
from sklearn.metrics import r2_score
```

```
Poly_Mae_train=metrics.mean_absolute_error(y_train,y_pred)
```

```
Poly_r2_train=r2_score(y_train,y_pred)
```

```
print('Error medio Absoluto (MAE):', metrics.mean_absolute_error(y_train, y_pred))
```

```
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_train, y_pred))
```

```
print('r2_score',r2_score(y_train, y_pred ))
```

```
Error medio Absoluto (MAE): 120111.24464567527
```

```
Root Mean Squared Error: 174164.006294397
```

```
r2_score 0.7706142422681812
```

```
polyX = poly_transform.fit_transform(X_test)
```

```
y_pred = poly_regression.predict(polyX)
```



```

from sklearn import metrics
from sklearn.metrics import r2_score
Poly_Mae=metrics.mean_absolute_error(y_test,y_pred)
Poly_r2=r2_score(y_test,y_pred)
print('Error medio Absoluto (MAE):', metrics.mean_absolute_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
print('r2_score',r2_score(y_test, y_pred ))

```

```

Error medio Absoluto (MAE): 127640.87573988685
Root Mean Squared Error: 212376.50824233124
r2_score 0.7137146909613681

```

ridge

```
X_train2=X_train.to_numpy().reshape(-1,1)
```

```

from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
model = Pipeline([("poly_features", PolynomialFeatures(degree=2, include_bias=False))
                  ("scaler", StandardScaler()),
                  ("linear_reg", LinearRegression())])
model.fit(X_train, y_train)

```

```

Pipeline(steps=[('poly_features', PolynomialFeatures(include_bias=False)),
                ('scaler', StandardScaler()),
                ('linear_reg', LinearRegression())])

```

```

from sklearn import metrics
from sklearn.metrics import r2_score
Ridge_Mae=metrics.mean_absolute_error(y_test,y_pred)
Ridge_r2=r2_score(y_test,y_pred)
print('Error medio Absoluto (MAE):', metrics.mean_absolute_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
print('r2_score',r2_score(y_test, y_pred ))

```

```

Error medio Absoluto (MAE): 127640.87573988685
Root Mean Squared Error: 212376.50824233124
r2_score 0.7137146909613681

```

lasso

```
from sklearn.linear_model import Lasso
```

```
model_lasso = Pipeline([("poly_features", PolynomialFeatures(degree=5, include_bias=False)),
                        ("scaler", StandardScaler()),
                        ("lasso", Lasso(alpha = 1, random_state = 42, tol=0.1, max_iter=10000))])

model_lasso.fit(X_train, y_train)

Pipeline(steps=[('poly_features',
                 PolynomialFeatures(degree=5, include_bias=False)),
               ('scaler', StandardScaler()),
               ('lasso', Lasso(alpha=1, random_state=42, tol=0.1))])

y_pred = model_lasso.predict(X_train)
from sklearn import metrics
from sklearn.metrics import r2_score

Lasso_MAE_train = metrics.mean_absolute_error(y_train, y_pred)
Lasso_r2_train = r2_score(y_train, y_pred)
print('Error medio Absoluto (MAE):', metrics.mean_absolute_error(y_train, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_train, y_pred)))
print('r2_score', r2_score(y_train, y_pred))

Error medio Absoluto (MAE): 113884.60002386104
Root Mean Squared Error: 162623.202988018
r2_score 0.8000070629856829

y_pred = model_lasso.predict(X_test)

from sklearn import metrics
from sklearn.metrics import r2_score
Lasso_MAE = metrics.mean_absolute_error(y_test, y_pred)
Lasso_r2 = r2_score(y_test, y_pred)
print('Error medio Absoluto (MAE):', metrics.mean_absolute_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
print('r2_score', r2_score(y_test, y_pred))

Error medio Absoluto (MAE): 127385.85985379248
Root Mean Squared Error: 254950.23977556612
r2_score 0.5874306511507208

Metricas = list()
Metricas.append('Lin train')
Metricas.append('Lin')
Metricas.append('Poly train')
Metricas.append('Poly')
Metricas.append('Ridge train')
Metricas.append('Ridge')
```

```

Metricas.append('Lasso train')
Metricas.append('Lasso')
MAE_list = list()
MAE_list.append(linear_Mae_train)
MAE_list.append(linear_Mae)
MAE_list.append(Poly_Mae_train)
MAE_list.append(Poly_Mae)
MAE_list.append(Ridge_Mae_train)
MAE_list.append(Ridge_Mae)
MAE_list.append(Lasso_MAE_train)
MAE_list.append(Lasso_MAE)

```

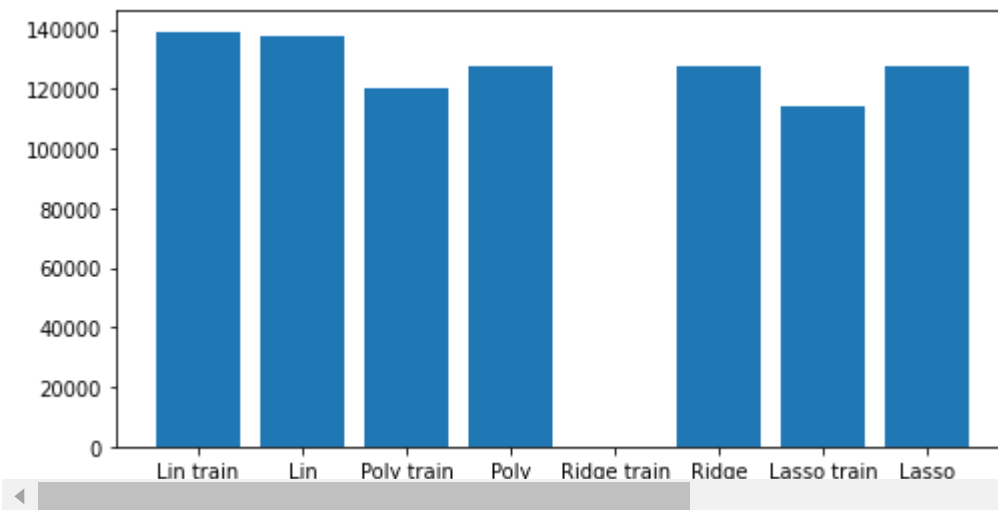
```

fig, ax = plt.subplots(figsize=(8, 4))
print("MAE")
print(MAE_list)
Metricas=list(Metricas)
MAE=list(MAE_list)
plt.bar(Metricas,MAE_list)

```

```
plt.show()
```

MAE  
[139269.32939115583, 137480.1388273178, 120111.24464567527, 127640.87573988685,

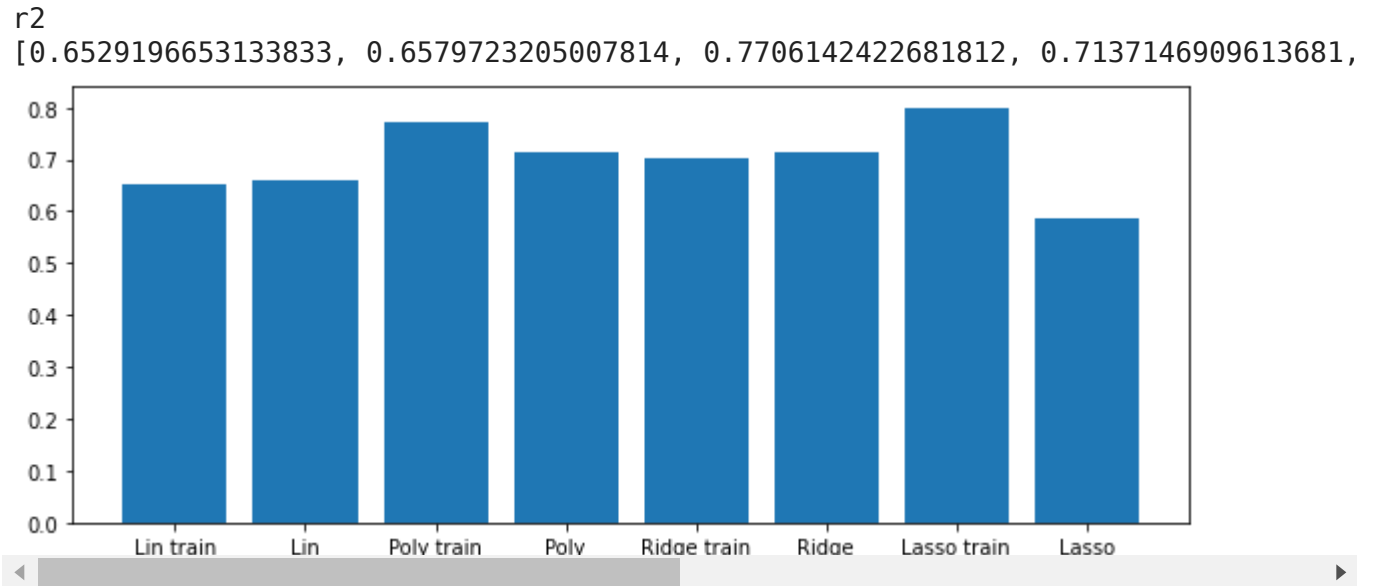


```

R2_list = list()
R2_list.append(linear_r2_train)
R2_list.append(linear_r2)
R2_list.append(Poly_r2_train)
R2_list.append(Poly_r2)
R2_list.append(Ridge_r2_train)
R2_list.append(Ridge_r2)
R2_list.append(Lasso_r2_train)
R2_list.append(Lasso_r2)
fig, ax = plt.subplots(figsize=(10, 4))
print("r2")
print(R2_list)

```

```
plt.bar(Metricas, R2_list)  
plt.show()
```



## Conclusiones

El comportamiento de Ridge es mas estable en cuanto a sus resultados. Pudiera ser una buena opcion. Dentro de todos los modelos, se vio que tuvo mejor desempeno y fue el mas estable.